



The time it takes to complete a query for a certain size grid is plotted in the graph above. These points are averages of 20 iterations of completing a random query within a certain grid size. All parts of the processing, including creating the grid, were included in every iteration. There was also a warmup period before collecting data where there each version was run 30 times. It seems like Version 5 is worse for larger grid sizes than Version 4, as Version 5 seems to have a quadratic line representing the time it will take to complete a query as the squares get larger, while Version 4 seems to only be linear. This seems to make sense as version 4 makes use of the ForkJoin framework, which is better able to portion out large tasks and fully use all resources, while version 5 is bound by only using 4 threads.